

BigData, i.e., Scalable Algorithm Design

The “Map Reduce” Programming Model

Marco Milanesio

UCA - MSc DATA SCIENCE 2019

- Jimmy Lin and Chris Dyer, “Data-Intensive Text Processing with MapReduce,” Morgan & Claypool Publishers, 2010¹
- Tom White, “Hadoop, The Definitive Guide,” O’Reilly / Yahoo Press, 2012
- Anand Rajaraman, Jeffrey D. Ullman, Jure Leskovec, “Mining of Massive Datasets”, Cambridge University Press, 2013
- Holden Karau, Andy Konwinski, Patrick Wendell and Matei Zaharia, “Learning Spark”, O’Reilly

This lecture is built starting from material from Prof. Michiardi’s “Cloud” course @Eurecom

¹<http://lintool.github.io/MapReduceAlgorithms/>

What is Big Data?

- **Vast repositories of data**

- The Web
- Physics
- Astronomy
- Finance

- **Volume, Velocity, Variety**

- **It's not the algorithm, it's the data!**

- More data leads to better accuracy
- With more data, accuracy of different algorithms converges

What is the “Map Reduce” Programming Model?

- **A distributed programming model:**
 - Inspired by functional programming
 - Inspired by Bulk Synchronous Parallelism (BSP)
- **An instance of an execution framework:**
 - Designed for large-scale data processing
 - Designed to run on clusters of commodity hardware

Key Principles

Scale out, not up!

- **For data-intensive workloads, a large number of commodity servers is preferred over a small number of high-end servers**
 - Cost of super-computers is not linear
 - But datacenter efficiency is a difficult problem to solve
- **Some numbers (\sim 2012):**
 - Data stored/processed by Google every day: $O(EB)$
 - Data stored/processed by Facebook every day: $O(PB)$

Implications of Scaling Out

- **Processing data is quick, I/O is very slow**
 - 1 Mechanical HDD \sim 100 MB/sec
 - 1000 Mechanical HDDs \sim 100 GB/sec
- **Sharing vs. Shared nothing:**
 - Sharing: manage a common/global state
 - Shared nothing: **independent** entities, no common state
- **Sharing is difficult:**
 - Synchronization, deadlocks
 - Finite bandwidth to access data from SAN
 - Temporal dependencies are complicated (restarts)

Failures are the norm, not the exception

- **Failures are part of everyday life**
 - Mostly due to the scale and shared environment
- **Sources of Failures**
 - Hardware / Software
 - Electrical, Cooling, ...
 - Unavailability of a resource due to overload
- **Failure Types**
 - Permanent
 - Transient

Move Processing to the Data

- **Drastic departure from high-performance computing model**
 - HPC: distinction between processing nodes and storage nodes
 - HPC: CPU intensive tasks
- **Data intensive workloads**
 - Generally not processor demanding
 - The network becomes the bottleneck
 - Framework generally assumes processing and storage nodes to be collocated

→ **Data Locality Principle**
- **Distributed filesystems are necessary**

Process Data Sequentially and Avoid Random Access

- **Data intensive workloads**

- Relevant datasets are too large to fit in memory
- Such data resides on disks

- **Disk performance is a bottleneck**

- **Seek times** for random disk access are **the problem**
 - Example: 1 TB DB with 10^{10} 100-byte records. Updates on 1% requires 1 month, reading and rewriting the whole DB would take 1 day²
- Organize computation for sequential reads

²From a post by Ted Dunning on the Hadoop mailing list

Implications of Data Access Patterns

- **Systems designed for:**
 - **Batch processing**
 - involving (mostly) **full scans** of the data
- **Typically, data is collected “elsewhere” and copied to the distributed filesystem**
 - E.g.: Apache Kafka, Hadoop Sqoop, ...
- **Data-intensive applications**
 - Read and process the whole Web (e.g. PageRank)
 - Read and process the whole Social Graph (e.g. LinkPrediction, a.k.a. “friend suggest”)
 - Log analysis (e.g. Network traces, Smart-meter data, ...)

Hide System-level Details

- **Separate the *what* from the *how***
 - Framework abstracts away the “distributed” part of the system
 - Such details are handled by internal primitives
- **BUT: In-depth knowledge of the framework is key**
 - Custom data reader/writer
 - Custom **data partitioning**
 - Memory utilization
- **Auxiliary components**
 - Too many to list!

Seamless Scalability

- **We can define scalability along two dimensions**
 - In terms of data: given twice the amount of data, the same algorithm should take no more than twice as long to run
 - In terms of resources: given a cluster twice the size, the same algorithm should take no more than half as long to run
- **Embarrassingly parallel problems**
 - Simple definition: independent (**shared nothing**) computations on fragments of the dataset
 - How to decide if a problem is embarrassingly parallel or not?

The Programming Model

Functional Programming Roots

- **Key feature: higher order functions**
 - Functions that accept other functions as arguments
 - **Map** and **Fold**

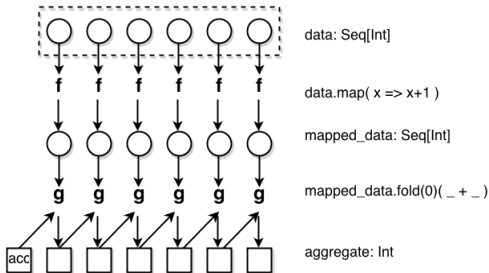


Figure: Illustration of *map* and *fold*.

Functional Programming Roots

- **map phase:**

- Given a list, *map* takes as an argument a function f (that takes a single argument) and applies it to all element in a list

- **fold phase:**

- Given a list, *fold* takes as arguments a function g (that takes two arguments) and an initial value (an accumulator)
- g is first applied to the initial value and the first item in the list
- The result is stored in an intermediate variable, which is used as an input together with the next item to a second application of g
- The process is repeated until all items in the list have been consumed

Functional Programming Roots

- **We can view map as a transformation over a dataset**
 - This transformation is specified by the function f
 - Each functional application happens in **isolation**
 - The application of f to each element of a dataset can be parallelized in a straightforward manner
- **We can view fold as an aggregation operation**
 - The aggregation is defined by the function g
 - Data locality: elements in the list must be “brought together”
 - If we can **group** elements of the list, also the fold phase can proceed in parallel
- **Associative and commutative operations**
 - Allow performance gains through local aggregation and reordering

Functional Programming and “Map Reduce”

- **Equivalence of “Map Reduce” and Functional Programming:**
 - The **map** of Hadoop MapReduce corresponds to the map operation
 - The **reduce** of Hadoop MapReduce corresponds to the fold operation
- **The framework coordinates the map and reduce phases:**
 - Grouping intermediate results happens in parallel
- **In practice:**
 - User-specified computation is applied (in parallel) to all input records of a dataset
 - Intermediate results are aggregated by another user-specified computation

What can we do with this Programming Model??

- **Introducing the Data Flow abstraction**

- The “old” Hadoop MapReduce programming model appears quite limited and strict
- Apache Spark programming model is much more flexible, and operates on a directed acyclic graph representative of the computations

- **Generally, everything can be computed with the “Map Reduce” model**

- We will focus on illustrative cases
- “design patterns”

Data Structures

- **Key-value pairs are the basic data structure in “Map Reduce”**
 - Keys and values can be: integers, float, strings, raw bytes
 - They can also be **arbitrary data structures**
- **The design of “Map Reduce” algorithms involves:**
 - Imposing the key-value structure on arbitrary datasets³
 - E.g.: for a collection of Web pages, input keys may be URLs and values may be the HTML content
 - In some algorithms, input keys are not used, in others they uniquely identify a record
 - Keys can be combined in complex ways to design various algorithms

³There's more about it: here we only look at the input to the map function.

A Generic “Map Reduce” Algorithm

- The programmer defines a mapper and a reducer as follows⁴⁵:
 - map: $(k_1, v_1) \rightarrow [(k_2, v_2)]$
 - reduce: $(k_2, [v_2]) \rightarrow [(k_3, v_3)]$
- In words:
 - A dataset stored on an underlying **distributed** filesystem, which is split in a number of **blocks** across machines
 - The mapper is applied to every input key-value pair to generate intermediate key-value pairs
 - The reducer is applied to all values associated with the same intermediate key to generate output key-value pairs

⁴We use the convention $[\dots]$ to denote a list.

⁵Pedices indicate different data types.

Where the magic happens

- **Implicit between the map and reduce phases is a **parallel “group by”** operation on intermediate keys**
 - Intermediate data arrive at each reducer in order, sorted by the key
 - No ordering is guaranteed across reducers
- **Output keys from reducers are written back to the distributed filesystem⁶**
 - The output may consist of r distinct files, where r is the number of reducers
 - Such output may be the input to a subsequent phase⁷

⁶Only Hadoop MapReduce. Apache Spark keeps in memory intermediate data.

⁷Think of **iterative algorithms**.

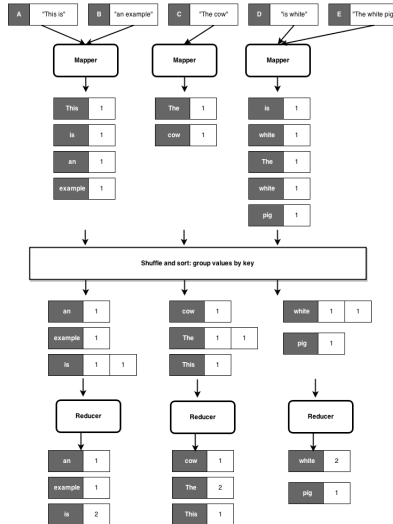
Where the magic happens

- **Intermediate keys are transient:**
 - They are not stored on the distributed filesystem
 - They are “spilled” to the local disk of each machine in the cluster

“Hello World” in “Map Reduce”

```
1: class MAPPER
2:   method MAP(offset  $a$ , line  $l$ )
3:     for all term  $t \in$  line  $l$  do
4:       EMIT(term  $t$ , count 1)

1: class REDUCER
2:   method REDUCE(term  $t$ , counts [ $c_1, c_2, \dots$ ])
3:      $sum \leftarrow 0$ 
4:     for all count  $c \in$  counts [ $c_1, c_2, \dots$ ] do
5:        $sum \leftarrow sum + c$ 
6:     EMIT(term  $t$ , count  $sum$ )
```



“Hello World” in “Map Reduce”

- **Input:**

- Key-value pairs: (offset, line) of a file stored on the distributed filesystem
- a: unique identifier of a line offset
- l: is the text of the line itself

- **Mapper:**

- Takes an input key-value pair, tokenize the line
- Emits intermediate key-value pairs: the word is the key and the integer is the value

- **The framework:**

- Guarantees all values associated with the same key (the word) are brought to the same reducer

- **The reducer:**

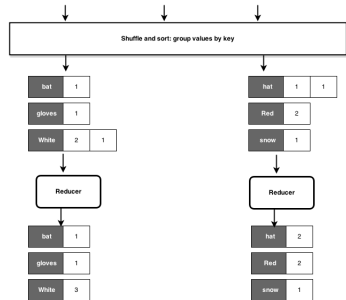
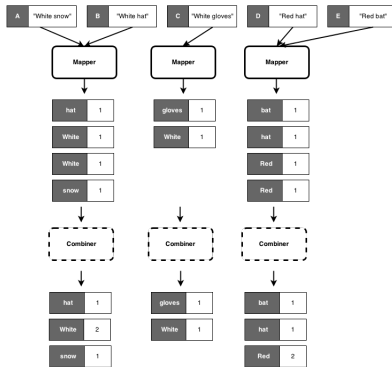
- Receives all values associated to some keys
- Sums the values and writes output key-value pairs: the key is the word and the value is the number of occurrences

Combiners

- **Combiners are a general mechanism to reduce the amount of intermediate data**
 - They could be thought of as “mini-reducers”
- **Back to our running example: word count**
 - Combiners aggregate term counts across documents processed by each map task
 - If combiners take advantage of all opportunities for local aggregation we have at most $m \times V$ intermediate key-value pairs
 - m : number of mappers
 - V : number of unique terms in the collection
 - Note: due to Zipfian nature of term distributions, not all mappers will see all terms

A word of caution

- **The use of combiners must be thought carefully**
 - In Hadoop, they are optional: the correctness of the algorithm cannot depend on computation (or even execution) of the combiners
 - In Apache Spark, they're mostly automatic
- **Combiners I/O types**
 - Input: $(k_2, [v_2])$ [Same input as for Reducers]
 - Output: $[(k_2, v_2)]$ [Same output as for Mappers]
- **Commutative and Associative computations**
 - Reducer and Combiner code may be interchangeable (e.g. Word Count)
 - This is not true in the general case



Algorithmic Correctness: an Example

- **Problem statement**

- We have a large dataset where input keys are strings and input values are integers
- We wish to compute the mean of all integers associated with the same key
 - In practice: the dataset can be a log from a website, where the keys are user IDs and values are some measure of activity

- **Next, a baseline approach**

- We use an **identity mapper**, which groups and sorts appropriately input key-value pairs
- Reducers keep track of running sum and the number of integers encountered
- The mean is emitted as the output of the reducer, with the input string as the key

Example: Computing the mean

```
1: class MAPPER
2:   method MAP(string  $t$ , integer  $r$ )
3:     EMIT(string  $t$ , integer  $r$ )

1: class REDUCER
2:   method REDUCE(string  $t$ , integers  $[r_1, r_2, \dots]$ )
3:      $sum \leftarrow 0$ 
4:      $cnt \leftarrow 0$ 
5:     for all integer  $r \in$  integers  $[r_1, r_2, \dots]$  do
6:        $sum \leftarrow sum + r$ 
7:        $cnt \leftarrow cnt + 1$ 
8:      $r_{avg} \leftarrow sum / cnt$ 
9:     EMIT(string  $t$ , integer  $r_{avg}$ )
```

Algorithmic Correctness

- **Note: operations are not distributive**

- $\text{Mean}(1,2,3,4,5) \neq \text{Mean}(\text{Mean}(1,2), \text{Mean}(3,4,5))$
- Hence: a combiner cannot output partial means and hope that the reducer will compute the correct final mean

- **Rule of thumb:**

- Combiners are optimizations, the algorithm should work even when “removing” them

Example: Computing the mean with combiners

```
1: class MAPPER
2:   method MAP(string t, integer r)
3:     EMIT(string t, pair (r, 1))
1: class COMBINER
2:   method COMBINE(string t, pairs [(s1, c1), (s2, c2) . . .])
3:     sum ← 0
4:     cnt ← 0
5:     for all pair (s, c) ∈ pairs [(s1, c1), (s2, c2) . . .] do
6:       sum ← sum + s
7:       cnt ← cnt + c
8:     EMIT(string t, pair (sum, cnt))
1: class REDUCER
2:   method REDUCE(string t, pairs [(s1, c1), (s2, c2) . . .])
3:     sum ← 0
4:     cnt ← 0
5:     for all pair (s, c) ∈ pairs [(s1, c1), (s2, c2) . . .] do
6:       sum ← sum + s
7:       cnt ← cnt + c
8:     ravg ← sum / cnt
9:     EMIT(string t, integer ravg)
```

Algorithm design

Algorithm Design

- **Developing algorithms involve:**
 - Preparing the input data
 - Implement the mapper and the reducer
 - Optionally, design the combiner and the partitioner
- **How to recast existing algorithms in “Map Reduce”?**
 - It is not always obvious how to express algorithms
 - Data structures play an important role
 - Optimization is hard
- **Learn by examples**
 - “Design patterns”
 - “Shuffle” is perhaps the most tricky aspect

Algorithm Design

- **Aspects that are *not* under the control of the designer**
 - *Where* a mapper or reducer will run
 - *When* a mapper or reducer begins or finishes
 - *Which* input key-value pairs are processed by a specific mapper
 - *Which* intermediate key-value pairs are processed by a specific reducer

Algorithm Design

- **Aspects that can be controlled**

- Construct **data structures as keys and values**
- Execute user-specified initialization and termination code for mappers and reducers
- Preserve state across multiple input and intermediate keys in mappers and reducers
- **Control the sort order** of intermediate keys, and therefore the order in which a reducer will encounter particular keys
- **Control the partitioning of the key space**, and therefore the set of keys that will be encountered by a particular reducer

Conclusions...

- **“Map Reduce” algorithms can be complex**
 - Hadoop MapReduce requires algorithm decomposition in several jobs
 - **Apache Spark is much simpler**
 - In general, iterative algorithms require a **driver**
 - Design patterns: http://www.dcs.bbk.ac.uk/~dell/teaching/cc/book/ditp/ditp_ch3.pdf